

Week 2: The Process abstraction

TABLE OF CONTENTS

1	Weekly Summary and Where are we?
a	Topics
b	Assignments
c	Reading Summary
2	Processes
a	Process provides each program with two key abstractions:
b	The illusion of <i>Multiprocessing</i>
c	Context switching
d	Modern systems are <i>multicore</i>
3	Error Handling in Unix/Linux
4	Working with Processes
a	Process IDs
b	Terminating processes
c	Creating processes
d	fork examples
e	Reaping child processes
f	Zombie Example
g	Non-Terminating Child Example
h	wait : Synchronizing with Children
i	Synchronizing with children
j	Another wait example
k	Summary
5	Building a Shell
a	A very simple shell
b	Adding output redirection and backgrounding.
c	Adding piping
d	Putting it all together
e	Why is this interesting?
f	Summary

Weekly Summary and Where are we?

Topics

- **Last Week:** Introduction to systems in general
- **This Week:** Introduction to the Process abstraction and how to use it in C
- **Next Week:** How the OS schedules processes (and other things)

Assignments

- **Lab 0** (setting up your dev environment) should be submitted already!
- **Lab 1** (getting familiar with C, `gdb`, `vim`, `ctags`, etc) is due on Friday
- **Lab 2** (building `ls`) is out; due 30 Jan 2026.
- **HW 1** is due on Friday

Reading Summary

- You should have caught up on Chap 1 and Chap 2 (Introduction to Operating Systems) from OSTEP last week
- Reading for this week was Chap 4 (the Process Abstraction) and Chap 5 (the Process API in C) from OSTEP
- **Chap 4** summarizes the need for a process abstraction
 - A **process** is an abstraction for a program that can be run
 - Once a process can be abstracted, the OS can run different processes at different times to most efficiently use the underlying hardware

- This is the first “thing” we’re virtualizing
- A process abstraction contains its state; all the things we need to know about it running
 - What lives in each register
 - process ID
 - parent process
 - What lives in each memory location
 - ...
- The process abstraction lives as an actual data structure that is provided
 - This is part of the API: Application Programming Interface that allows anyone to work with processes as long as they use the agreed-upon structures and format
- **Chap 5** introduces the Process API for Unix
 - `fork()`: Runs a new instance of the current program
 - `wait()`: Tells this process
 - `exec()`

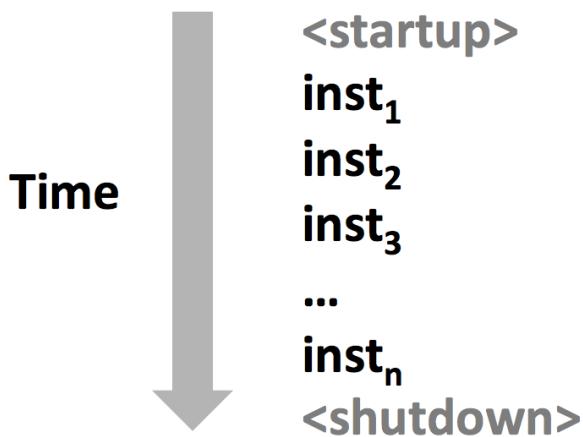
Processes

Definition: A process is an instance of a running program.

- One of the most profound ideas in computer science
- Not the same as “program” or “processor”

A processor ONLY reads a sequence of instructions from start to finish; this is the ***control flow***.

Physical control flow



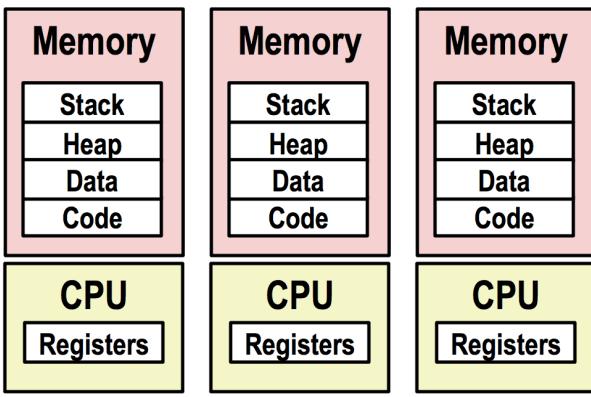
A processor runs a process.

Process provides each program with two key abstractions:

- ***Logical control flow***
 - Each program seems to have exclusive use of the CPU
 - Provided by kernel mechanism called context switching
- ***Private address space***
 - Each program seems to have exclusive use of main memory
 - Provided by kernel mechanism called ***virtual memory***

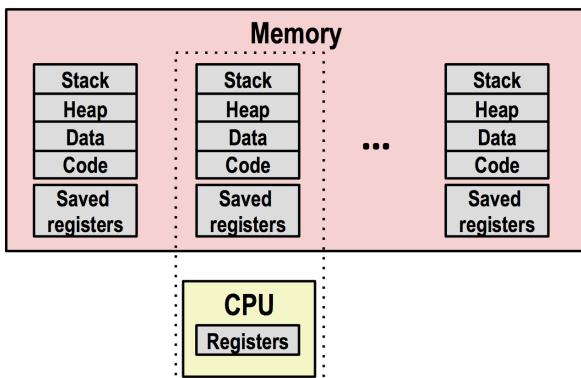
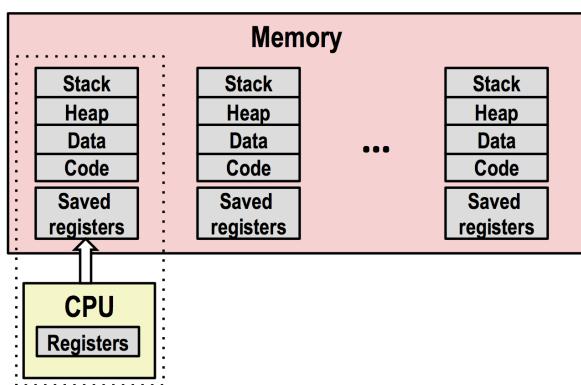
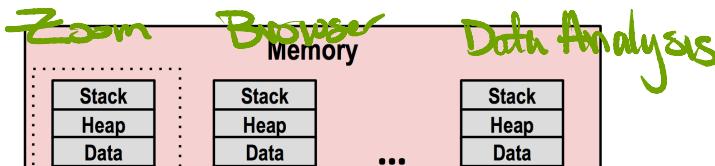
The illusion of *Multiprocessing*

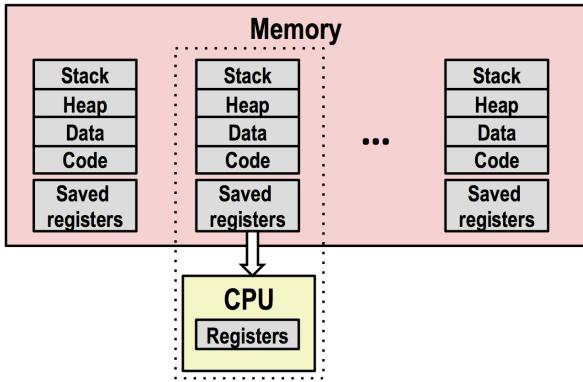
- See [top](#)



The reality, traditionally:

- Computers have one processor that gives the illusion of being able to run multiple programs/processes at once
- Single processor executes multiple processes concurrently
 - Process executions interleaved (multitasking)
 - Address spaces managed by virtual memory system (later in course)
 - Register values for nonexecuting processes saved in memory

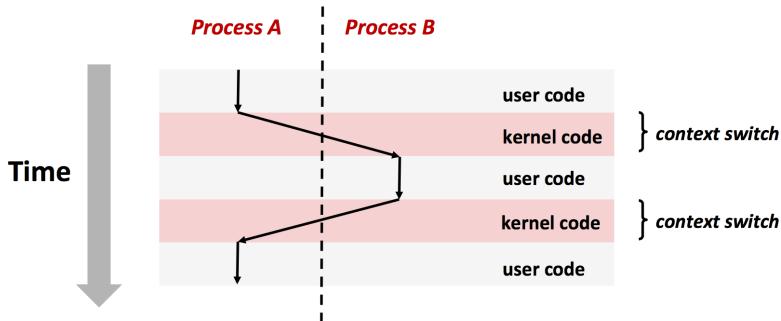




Context switching

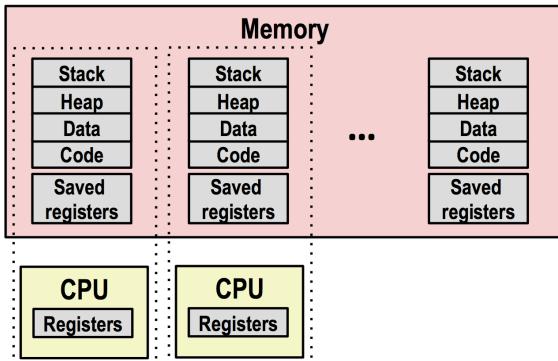
Processes are managed by a shared chunk of memory resident OS code called the kernel

- Important: the kernel is not a separate process, but rather runs as part of some existing process.
- Control flow passes from one process to another via a ***context switch***



Modern systems are *multicore*

- Machines have multiple processors
 - Multiple CPUs on single chip
 - Share main memory (and some of the caches)
 - Each can execute a separate process
 - Scheduling of processors onto cores done by ***kernel***



Error Handling in Unix/Linux

On error, Linux system-level functions typically return `-1` and set global variable `errno` to indicate cause.

- *Hard and fast rule:*
- You must check the return status of every system-level function
- Only exception is for the few functions that return void
- Example:

```
if ((pid = fork()) < 0) {
    fprintf(stderr, "fork error: %s\n", strerror(errno));
    exit(0);
}
```

Can utilize an error-reporting function:

```
void unix_error(char *msg) /* Unix-style error */
{
    fprintf(stderr, "%s: %s\n", msg, strerror(errno));
    exit(0);
}

if ((pid = fork()) < 0)
    unix_error("fork error");
```

Can simplify more by wrapping system calls with an error-handling function:

```
pid_t Fork(void)
{
    pid_t pid;
    if ((pid = fork()) < 0)
        unix_error("Fork error");
    return pid;
}

pid = Fork();
```

Working with Processes

Things we want to do with processes

- Get process IDs
- Create and terminate processes
- Reap child processes
- Kill the Zombies!
- Synchronizing with child processes

Process IDs

- `pid_t getpid(void)`
 - Returns PID of current process
- `pid_t getppid(void)`
 - Returns PID of parent process

From a programmer's perspective, we can think of a process as being in one of three states:

- **Running**
 - Process is either executing, or waiting to be executed and will eventually be scheduled (i.e., chosen to execute) by the kernel
- **Stopped**

- Process execution is suspended and will not be scheduled until further notice (future class we'll study signals)
- **Terminated**
- Process is stopped permanently

Terminating processes

- Process becomes terminated for one of three reasons:
 - Receiving a signal whose default action is to terminate (future class)
 - Returning from the `main` routine
 - Calling the `exit` function
- `void exit(int status)`
 - Terminates with an exit status of `status`
 - Convention: normal return status is 0, nonzero on error
 - Another way to explicitly set the exit status is to return an integer value from the main routine
- `exit` is called once but never returns.

Creating processes

- Parent process creates a new running child process by calling `fork`
- `int fork(void)`
 - Returns 0 to the child process, child's PID to parent process
 - Child is *almost* identical to parent:
 - Child get an identical (but separate) copy of the parent's virtual address space.
 - Child gets identical copies of the parent's open file descriptors
 - Child has a different PID than the parent
- `fork` is interesting (and often confusing) because it is called once but returns twice

`fork` examples

```
int main()
{
    pid_t pid;
    int x = 1;

    pid = fork(); // Consider using our error-catching Fork() from above
    if (pid == 0) { /* Child */
        printf("child : x=%d\n", ++x);
        exit(0);
    }
    /* Parent */
    printf("parent: x=%d\n", --x);
    exit(0)
}
```

- Call once, return twice
- Concurrent execution
 - Can't predict execution order of parent and child
- Duplicate but separate address space
 - `x` has a value of 1 when `fork` returns in parent and child
 - Subsequent changes to `x` are independent
- Shared open files
 - `stdout` is the same in both parent and child

Reaping child processes

Idea:

- When process terminates, it still consumes system resources
 - Examples: Exit status, various OS tables
- Called a "zombie"
 - Living corpse, half alive and half dead
- Reaping
 - Performed by parent on terminated child (using `wait` or `waitpid`)
 - Parent is given exit status information
 - Kernel then deletes zombie child process
- What if parent doesn't reap?
 - If any parent terminates without reaping a child, then the orphaned child will be reaped by init process (`pid == 1`)
- So, only need explicit reaping in long-running processes
 - e.g., shells and servers

Zombie Example

- `ps` shows child process as "defunct" (i.e., a zombie)
- Killing parent allows child to be reaped by init

Non-Terminating Child Example

- `ps` shows child process as "defunct" (i.e., a zombie)
- Killing parent allows child to be reaped by init

`wait` : Synchronizing with Children

- Parent reaps a child by calling the `wait` function
- `int wait(int *child_status)`
 - Suspends current process until one of its children terminates
 - Return value is the pid of the child process that terminated
 - If child status != NULL, then the integer it points to will be set to a value that indicates reason the child terminated and the exit status:
 - Checked using macros defined in `wait.h`
 - `WIFEXITED`, `WEXITSTATUS`, `WIFSIGNALED`, `WTERMSIG`, `WIFSTOPPED`, `WSTOPSIG`, `WIFCONTINUED`
 - Details... look them up if you're interested :)

Synchronizing with children

```
void fork() {
    int child_status;
    if (fork() == 0) {
        printf("HC: hello from child\n");
        exit(0);
    } else {
        printf("HP: hello from parent\n");
        wait(&child_status);
        printf("CT: child has terminated\n");
    }
    printf("Bye\n");
}
```

Feasible output: HC → HP → CT → Bye

Infeasible output: HP → CT → Bye → HC

Another `wait` example

```
void fork_wait() {
    pid_t pid[N];
```

```

int i, child_status;
for (i = 0; i < N; i++) {
    if ((pid[i] = fork()) == 0) {
        exit(100+i); /* Child */
    }

    for (i = 0; i < N; i++) { /* Parent */
        pid_t wpid = wait(&child_status);
        if (WIFEXITED(child_status))
            printf("Child %d terminated with exit status %d\n",
                   wpid, WEXITSTATUS(child_status));
        else
            printf("Child %d terminate abnormally\n", wpid);
    }
}

```

- If multiple children completed, will take in arbitrary order
- Can use macros `WIFEXITED` and `WEXITSTATUS` to get information about exit status

Summary

- Processes are the abstraction that allows operating systems to utilize hardware resources efficiently
- Processes have a context, which is basically all their data and instructions, and where those items live in memory/hardware
- Before a context switch, the kernel will save a process context for the process currently being executed.
 - After the current process is stored, the next process is restored from its stored context
 - The new process continues its execution
- As programmers, we can make use of processes by using `fork()` and `exec()` in C.
- Once we start creating new processes, we need to worry about whether our instructions are being executed concurrently or sequentially

Building a Shell

When we speak about "the command line", we're actually referring to **the shell**. The *shell* is a program that takes keyboard commands and passes them to the operating system to execute. `bash` (Bourne Again Shell) is a common shell included with most Linux distributions; `zsh` (Z Shell) is a default shell available on Macs.

Now that we know how to work with processes via `fork()` and `exec()`, we can actually build our own shell at this point.

A very simple shell

```

while (1)
{
    Wait_Status
    write(1, "$ ", 2);
    readcommand(command, args) // parse input
    if ((pid = fork()) == 0)
    {
        // child
        execve(command, args, 0);

    } else if (pid > 0)
    {
        // parent
        wait(0); // wait for child
    } else
}

```

ls <cr>

```

    {
        perror("Failed to fork");
    }
}

```

Adding output redirection and backgrounding.

Output redirection is something like `ls > list.txt`, where we redirect the output of the `ls` program to the file `list.txt`. Backgrounding means something like `myprog &`, which allows the program to run in the background, not taking up the shell's prompt.

```

while(1)
{
    write(1, "$ ", 2);
    readcommand(command, args); // parse input
    if ((pid = fork()) == 0)
    {
        // child
        if (output_redirected)
        {
            close(1);
            open(redirect_file, O_CREAT | O_TRUNC | O_WRONLY, 0666);
        }
        // When the command runs, fd 1 will refer to the redirected file
        execve(command, args, 0);

    } else if (pid > 0)
    {
        // parent
        if (foreground_process)
        {
            wait(0); // wait for child
        }
    } else
    {
        perror("Failed to fork");
    }
}

```

Adding piping

Another example: the `pipe()` syscall. This allows something like `$ ls | sort | head -4`.

```

int fdarray[2];
char buf[512];
int n;

pipe(fdarray);
write(fdarray[1], "hello", 5);

n = read(fdarray[0], buf, sizeof(buf));
// buf[] now contains 'h', 'e', 'l', 'l', 'o'

```

File descriptors are inherited across `fork()`.

```
// How two processes can communicate over a pipe
int fdarray[2];
char buf[512];
int n, pid;

pipe(fdarray);
pid = fork();
if (pid > 0)
{ PARENT
    write(fdarray[1], "hello", 5);
}
else
{
    n = read(fdarray[0], buf, sizeof(buf));
}
* Does Read block ?
```

Putting it all together

Putting it all together: implement a shell using `fork()`, `exec()` and `pipe()`.

```
void main_loop()
{
    while(1)
    {
        write(1, "$ ", 2);
        readcommand(command, args); // parse input
        if ((pid = fork()) == 0)
        {
            // child
            if (pipeline_requested)
            {
                handle_pipeline(left_command, right_command);
            }
            else
            {
                if (output_redirected)
                {
                    close(1);
                    open(redirect_file, O_CREAT | O_TRUNC | O_WRONLY, 0666);
                }
                exec(command, args, 0);
            }
        } else if (pid > 0)
        {
            // parent
            if (foreground_process)
            {
                wait(0);
            }
        } else {
            perror("Failed to fork");
        }
    }
}
```

ls | sort

```

void handle_pipeline(left_command, right_command)
{
    int fdarray[2];
    if (pipe(fdarray) < 0) panic ("error");

    if ((pid = fork()) == 0)
    {
        // child; left command
        dup2(fdarray[1], 1); // make fd 1 the same as fdarray[1], which is the write end of the pipe
        // implies close(1)

        close(fdarray[0]);
        close(fdarray[1]);
        parse(command1, args1, left_command);
        exec(command1, args1, 0);
    }
    else
    {
        if (pid > 0)
        {
            // parent; right command
            dup2(fdarray[0]); // make fd 0 the same as fdarray[0], which is the read end of the pipe
            // implies close(0)

            close(fdarray[0]);
            close(fdarray[1]);
            parse(command2, args2, right_command);
            exec(command2, args2, 0);
        }
        else
        {
            printf("Unable to fork\n");
        }
    }
}
}

```

Why is this interesting?

- Pipelines and output redirection are done by manipulating the child's environment, not by asking a program author to implement a complex set of behaviors.
- The same code for `ls` can result in:
 - printing to the screen (`ls -l`)
 - writing to a file (`ls -l > output.txt`)
 - getting the output from `ls` formatted by another program (`ls -l | sort`)
- Without redirection, the `ls` implementation would have to anticipate every possible output mode and would build in an interface so the user (client) could specify how the output is treated
- Rather, the author of `ls` wrote it in terms of a file descriptor: `write(1, "some output", byte_count)`.
- Through the use of `fork()` and `exec()`, the shell can make that file descriptor be a pipe, a file, the console, etc.

Summary

- ***

- Process Abstraction
- fork(), waiting for child processes to terminate.
- Write your own exceptions
- Redirecting and piping
- Macros for exit status